Scholarix AI — Customer Interview Report

Two interviews conducted as part of the two-week team evaluation task

Interview 1 — Architecture Grad Student

Role	Architecture Graduate Student
Why relevant	Actively scoping research papers and literature — uses researcher profiles to understand what's been published and where their own work fits. Represents a core user type for Scholarix AI.
Job to be done	Find credible, up-to-date research to scope a paper and identify gaps in the literature.
Issue faced	When researching a paper's scope, relies on whatever is available across platforms without a clear signal of whether information is current or complete. Cross-checks across platforms because information can be missing or removed from individual sources.
Why they'd use this	A tool that confirms information from multiple sources would help them go deeper. They said they would 'go further into the multiple sources to find out more information.' The multi-source model matches how they already work.
Trust concerns	They currently trust sources they already consider reliable — but have no strong framework for evaluating unfamiliar ones. This means they may not know when data is wrong. A tool that surfaces source disagreement directly addresses this blind spot.
Features wanted	Multi-source confirmation, relevance filtering, cross-platform coverage. They would stop using a tool that returns irrelevant results.

Interview Questions — Rationale & Findings

Q1. Walk me through the last time you looked up a researcher or checked someone's publications — what were you trying to find out?

"I was researching on what is published and what information is out there and where could my research paper scope be."

Why this question was chosen: To get a real memory instead of a hypothetical. We needed to understand the actual job they were trying to do, not an abstract opinion.

What we learned: The user is scoping literature — they're not evaluating a specific researcher deeply, but mapping what exists. This means breadth and relevance matter more to them than deep verification.

How it affected the prototype: Reinforced that the tool needs to surface information clearly and relevantly. A complex audit interface would not fit this user — they need a scannable, clear output.

Q2. When you're looking at a researcher's profile, how do you decide whether to trust the information?

"I use reliable sources so I can trust the researcher's information."

Why this question was chosen: To understand their existing trust framework and whether they have one.

What we learned: They rely on the source's reputation rather than cross-checking the content itself. This means they don't currently detect errors — they assume reliable sources are accurate. This is the exact blind spot the tool addresses.

How it affected the prototype: Validated the need for source attribution labels. If a user assumes a source is reliable, showing them where information came from — and when sources conflict — challenges that assumption in a constructive way.

Q3. Have you ever found information about a researcher that turned out to be wrong or out of date? What happened?

"No."

Why this question was chosen: To surface real examples of data failure and understand the personal stakes.

What we learned: They have not personally encountered visible errors — which means errors are either not happening in their workflow or they're not catching them. Both are significant. The second possibility (not catching errors) is supported by Q2, where they said they trust source reputation without cross-checking.

How it affected the prototype: Reinforced the importance of making errors visible proactively, not just after users notice them. The tool needs to surface issues the user wouldn't find on their own.

Q4. If a tool showed you that some information was confirmed from multiple sources and some couldn't be verified — would that be useful, or would it create more confusion?

"It would definitely be more useful. I would further go into the multiple sources and find out more information. It would help me get more information."

Why this question was chosen: To directly test the core product idea — confidence labelling with source attribution.

What we learned: The user responded positively and described exactly the behaviour we want to enable: using the tool as a starting point to go deeper. This validates the 'Verified / Uncertain / Suspicious' label system.

How it affected the prototype: Confirmed that showing source breakdown is the right direction. The user will engage with it rather than be confused by it.

Q5. What would make you stop using a tool like this entirely?

"If it's not giving me relevant information."

Why this question was chosen: To identify hard limits — what would cause complete abandonment rather than reduced use.

What we learned: Relevance is the primary threshold. The user doesn't require perfection, but they require that results connect to what they're actually looking for.

How it affected the prototype: Set a hard requirement: every flagged issue must be explained in plain English. A flag without a clear reason looks like noise, not signal.

Q6. When looking for collaborators or doing a literature review, do you ever check the same researcher across multiple platforms? Why or why not?

"Yes, I check the same researcher across multiple platforms because some information might be taken down through other sources. Just to check if the information is fully accurate."

Why this question was chosen: To understand whether users are already doing manual cross-platform verification — which is exactly the painful process our tool automates.

What we learned: They already cross-check manually, driven by awareness that individual sources can be incomplete or removed. This directly validates the cross-platform comparison feature.

How it affected the prototype: Confirmed that the affiliation conflict checker and citation source comparison panel are solving a problem users are already trying to solve manually.

Interview 2 — CS Grad Student (RA)

Role	CS Graduate Student and Research Assistant — actively reading, evaluating, and citing papers for course projects and independent research in agentic AI and LLMs.
Why relevant	Represents a technically sophisticated research user who evaluates researcher profiles at depth — checking citations, author background, affiliation, GitHub reproducibility, and cross-platform agreement. This is the power user profile for Scholarix AI.
Job to be done	Evaluate whether a paper and its authors are credible enough to build on — checking citation count, affiliation, abstract alignment, and result reproducibility before committing time to reading the full paper.
Issue faced	Research data goes stale fast. In agentic AI, a paper's findings can be invalidated within months by a new model release. He's repeatedly found papers where the baseline comparisons are outdated. He also encounters

	author disambiguation failures — common surnames lead to misattributed papers, especially on Google Scholar.

Why they'd use this	Already performs multi-platform verification manually. Described Google Scholar as having 'the widest coverage but the worst hygiene — auto-attributed papers, name collisions.' Our tool automates what he does by hand.

Trust concerns	Two distinct trust concerns: (1) data staleness in fast-moving fields; (2) hallucination risk in AI-based tools — he said he would verify even confirmed information if the tool is LLM-based, because of grounding errors. He cross-checks lab pages to confirm whether a researcher is still active and taking students.

Features wanted	Citation count from multiple sources, author disambiguation, staleness flags, lab page affiliation check, GitHub/reproducibility link, clear source provenance so he can verify himself.

Interview Questions — Rationale & Findings

Q1. Walk me through the last time you looked up a researcher or checked someone's publications — what were you trying to find out?

"I was working on a course project about agentic AI. I found a paper called 'Get my drift.' First thing I did was check the author's backgrounds, citation and abstract. All of these give cues as to how popular it is within the community, and how does it connect to my work. I decided to read the paper, especially the results part, since they give extra validation about the depth of the study and whether all of my questions are answered."

Why this question was chosen: To understand the real workflow — what information they check, in what order, and what decision each piece of information informs.

What we learned: He uses a layered evaluation: author background → citation count → abstract alignment → results section. Each layer is a trust checkpoint. This is exactly the audit structure our prototype uses.

How it affected the prototype: The layered check sequence (profile → affiliation → publications → confidence scores) maps directly to his natural evaluation workflow. The prototype's structure mirrors how this user already thinks.

Q2. When you're looking at a researcher's profile, how do you decide whether to trust the information?

"I check the article's citation, the researcher's total citations, their Google Scholar account, and whether the paper has a GitHub repository link attached — since that adds validation that the authors have open-sourced the code. I'll try to run the repository once to make sure the results are reproducible."

Why this question was chosen: To understand their trust signals and whether they have a structured framework or rely on intuition.

What we learned: He has a structured, multi-signal trust framework: citations, Google Scholar, and reproducibility via GitHub. Notably, he already knows Google Scholar can be wrong — he uses it as one signal, not a definitive answer.

How it affected the prototype: Confirmed that showing citation counts from multiple sources side by side is the right approach. Also surfaced reproducibility as a trust signal we hadn't prioritised — worth noting as a future feature.

Q3. Have you ever found information about a researcher that turned out to be wrong or out of date? What happened?

"Yes, many times. In agentic AI and VLMs, it happens so often that a lab releases a model that performs much better on benchmarks, and since the study was conducted before the release, the paper's conclusions are already outdated. Sometimes the new frontier models themselves can solve the problem, defeating the purpose of the research."

Why this question was chosen: To get a concrete example of data failure and understand the real-world impact.

What we learned: Data staleness is a genuine, recurring problem in fast-moving fields. This isn't theoretical — it's something he experiences regularly. He frames it as a field-level problem, not just a tool problem.

How it affected the prototype: Validates the need for staleness flags. A paper published in 2022 about a capability that GPT-4 now handles trivially is misleading without context. Future versions of the tool should flag publication year relative to benchmark dates in the field.

Q4. If a tool showed you that some information was confirmed from multiple sources and some couldn't be verified — would that be useful, or would it create more confusion?

"That would be useful. Some information is verified and some couldn't be — so I would instantly go to work on verifying the unverified information. But I would also check the already-verified information, because if it is an LM-based tool there is a slight chance it might ground the responses incorrectly due to hallucination. So I would verify both ways, to err on the side of caution."

Why this question was chosen: To test the core product concept and understand how a sophisticated user would interact with confidence labels.

What we learned: He would use the tool as a starting point, not a final answer — which is exactly the right behaviour to design for. He also raised an important concern: if the tool itself is LLM-based, its verified claims still need checking. This is a critical trust design insight.

How it affected the prototype: Reinforced that the tool must show source provenance explicitly — not just a label, but where the label came from. 'Verified by OpenAlex and Crossref' is more trustworthy than just 'Verified.' We added source attribution to every confidence label.

Q5. What would make you stop using a tool like this entirely?

Implied by his other answers: if the tool presents confident labels without showing the source, or if it hallucinates verified status for information that isn't actually confirmed.

Why this question was chosen: To surface hard limits and design constraints.

What we learned: For this user, opacity is the deal-breaker. He needs to see where every claim came from so he can verify it himself. A black-box confidence score would not work for him.

How it affected the prototype: Made source attribution a non-negotiable feature, not a nice-to-have. Every flag in the prototype shows which source produced it.

Q6. When looking for collaborators or doing a literature review, do you ever check the same researcher across multiple platforms? Why or why not?

"Yes, mainly because no single source is both complete and correct. Google Scholar has the widest coverage but the worst hygiene — auto-attributed papers, name collisions that get bad with common surnames. DBLP disambiguates CS authors cleanly but still has its own issues. OpenReview reveals how a line of work developed through revisions. A lab page tells you whether they're still there and taking students. Cross-checking also lets me correct my own misattributions rather than propagate them."

Why this question was chosen: To understand manual verification behaviour and whether users are already solving the problem our tool automates.

What we learned: He is doing exactly what our tool automates — manually checking Google Scholar, DBLP, OpenReview, and lab pages. His description of Google Scholar's 'name collisions that get bad with common surnames' is a direct real-world parallel to the Boxuan Zhao disambiguation failure we found in the dataset.

How it affected the prototype: This single answer validated three features at once: multi-source citation comparison, author disambiguation flags, and affiliation verification against current lab pages. It also gave us language to use in the pitch: we are automating a process this user already does manually, every time.

Customer Discovery → Prototype Features

The table below maps specific interview findings to features built in the prototype. Every major feature can be traced to something a real user said.

What the user said	What we built
"No single source is both complete and correct" (Interviewee 2)	Citation source comparison panel — OpenAlex vs Crossref side by side
"Google Scholar has the worst hygiene — name collisions with common surnames" (Interviewee 2)	Author disambiguation flag — surfaces Boxuan Zhao-style wrong-person attribution
"I check the lab page to see if they're still there" (Interviewee 2)	ORCID-affiliation conflict checker — flags profile vs ORCID institution mismatch
"I check across multiple platforms because information might be taken down" (Interviewee 1)	Multi-source verification with MATCH / CONFLICT / CANNOT VERIFY labels
"I'd stop using it if it's not giving me relevant information" (Interviewee 1)	Plain-English reason shown for every flag — no unexplained noise
"I would verify even confirmed info if it's an LLM-based tool" (Interviewee 2)	Every confidence label shows the source it came from — no black-box scores

"I check whether there's a GitHub repo — that's a validation check" (Interviewee 2)	Noted as a future feature: reproducibility link presence check

Both interviewees are graduate students with direct experience searching for, evaluating, and citing researcher profiles and publications.

Web Application ScreenShots:

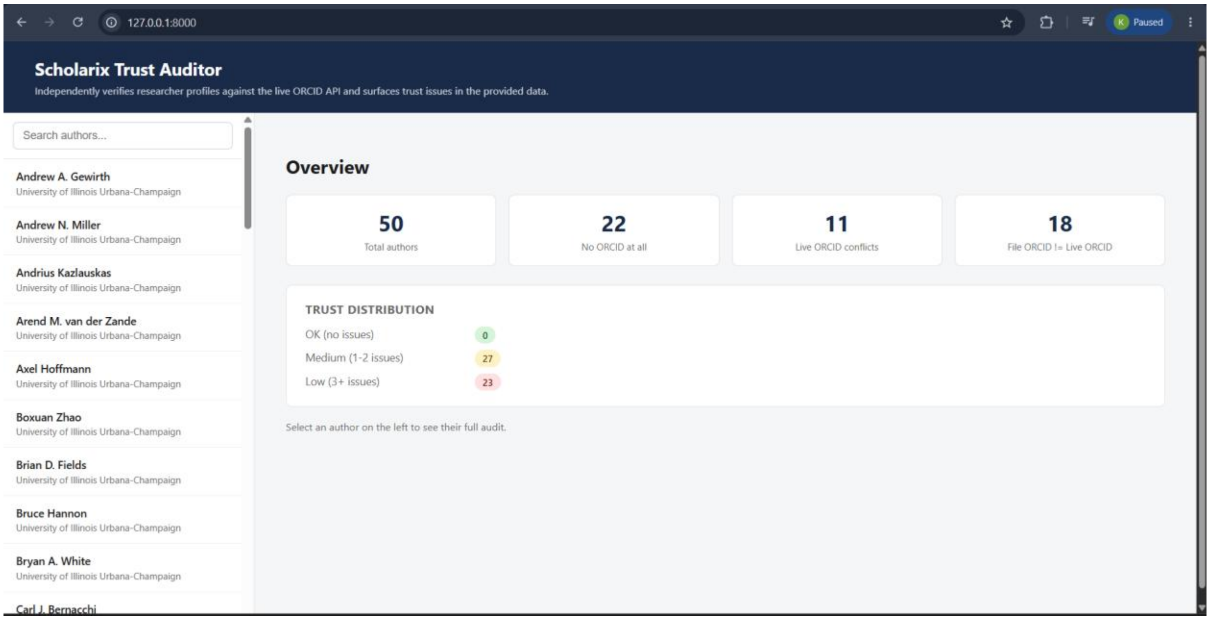

Low trust author:

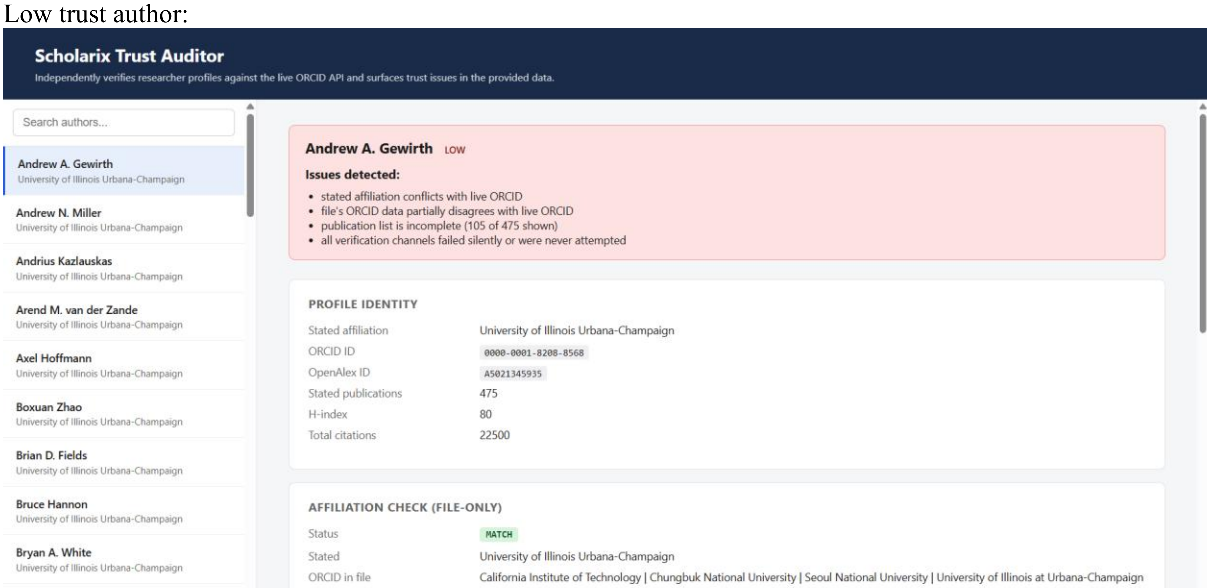

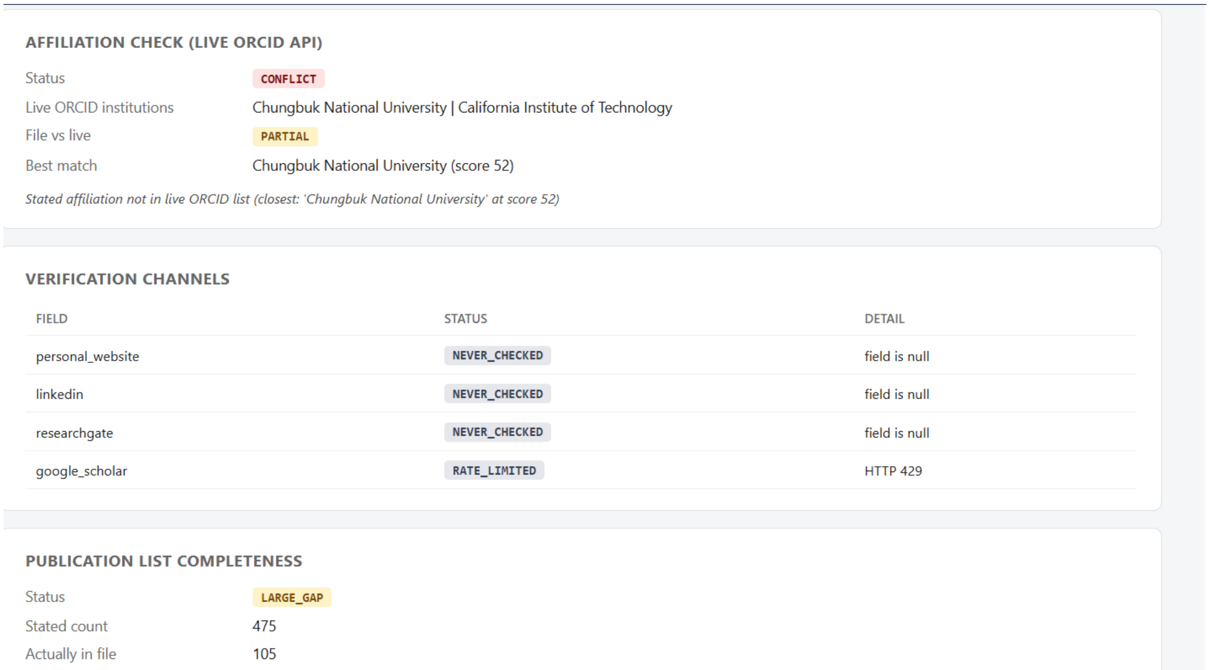

Medium trust author:

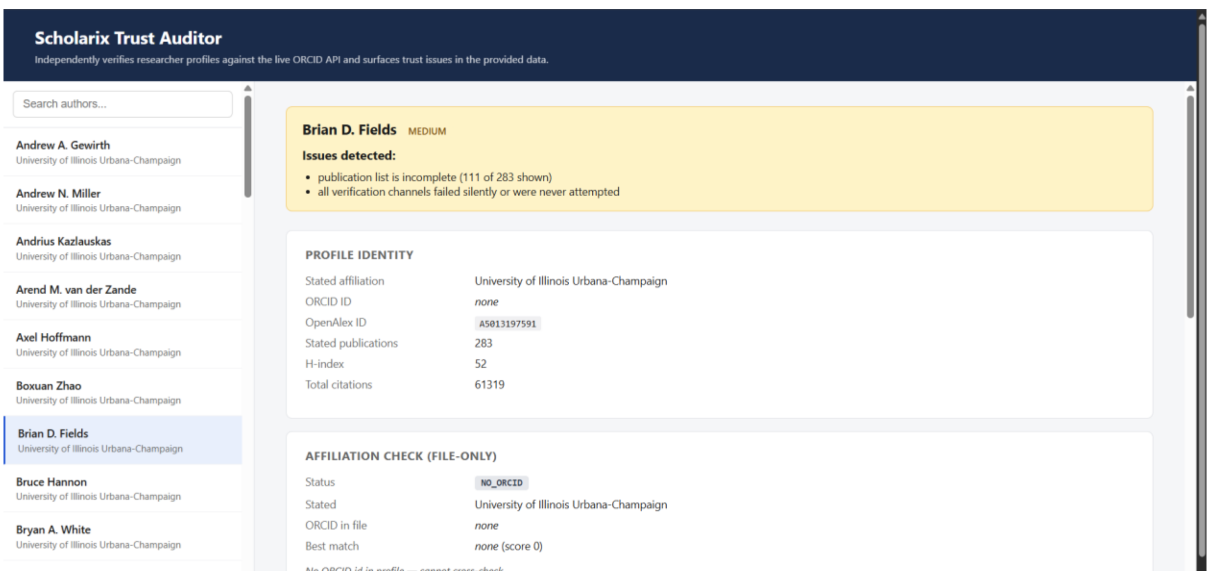

AFFILIATION CHECK (LIVE ORCID API)

Status	NO_ORCID
Live ORCID institutions	*none*
File vs live	*n/a*
Best match	*none* (score 0)

No ORCID id in profile

VERIFICATION CHANNELS

FIELD	STATUS	DETAIL
personal_website	NEVER_CHECKED	field is null
linkedin	NEVER_CHECKED	field is null
researchgate	NEVER_CHECKED	field is null
google_scholar	RATE_LIMITED	HTTP 429

PUBLICATION LIST COMPLETENESS

Status	LARGE_GAP
Stated count	283
Actually in file	111